

# CS330 - Computer Organization and Assembly Language Programming

Lecture 15

-Review -

Professor : Mahmut Unan – UAB CS

# Agenda

## Review

*You responsible for all the topics*

*This is a rough review*

## Midterm Exam

October 1<sup>st</sup>, 2021 Friday

Lecture Time

# Exam structure

- Multiple choice questions: ~20
- Conversions(binary,decimal,hex,mb,tb..etc), calculations (Amdahl's law)
- Open ended questions

# Hello World !

A typical *C program* basically consists of the following parts;

- Preprocessor Commands
- Functions
- Variables
- Statements & Expressions
- Comment

```
#include <stdio.h>
int main()
{
    /* the first program in CS330 */
    printf("hello, world\n");
    return 0;
}
```

# Information is Bits + Context

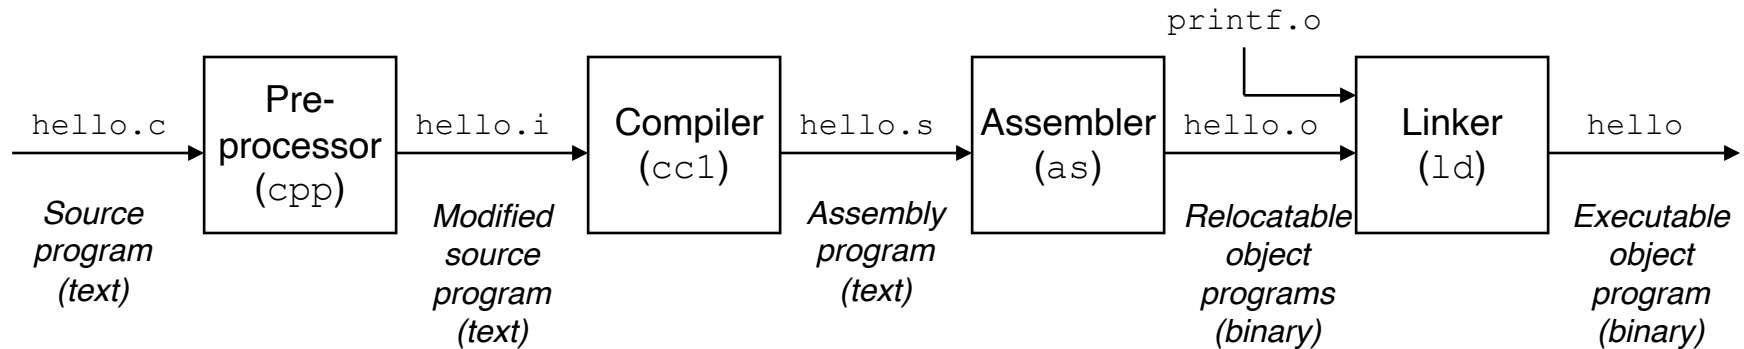
- “hello.c” is a source code
  - Sequence of bits (0 or 1)
  - 8-bit data chunks are called bytes
  - Each byte represents some text character in the program

#	i	n	c	l	u	d	e	<sp>	<	s	t	d	i	o	.
35	105	110	99	108	117	100	101	32	60	115	116	100	105	111	46
h	>	\n	\n	i	n	t	<sp>	m	a	i	n	(	)	\n	{
104	62	10	10	105	110	116	32	109	97	105	110	40	41	10	123
\n	<sp>	<sp>	<sp>	<sp>	p	r	i	n	t	f	(	"	h	e	l
10	32	32	32	32	112	114	105	110	116	102	40	34	104	101	108
l	o	,	<sp>	w	o	r	l	d	\	n	"	)	;	\n	}
108	111	44	32	119	111	114	108	100	92	110	34	41	59	10	125

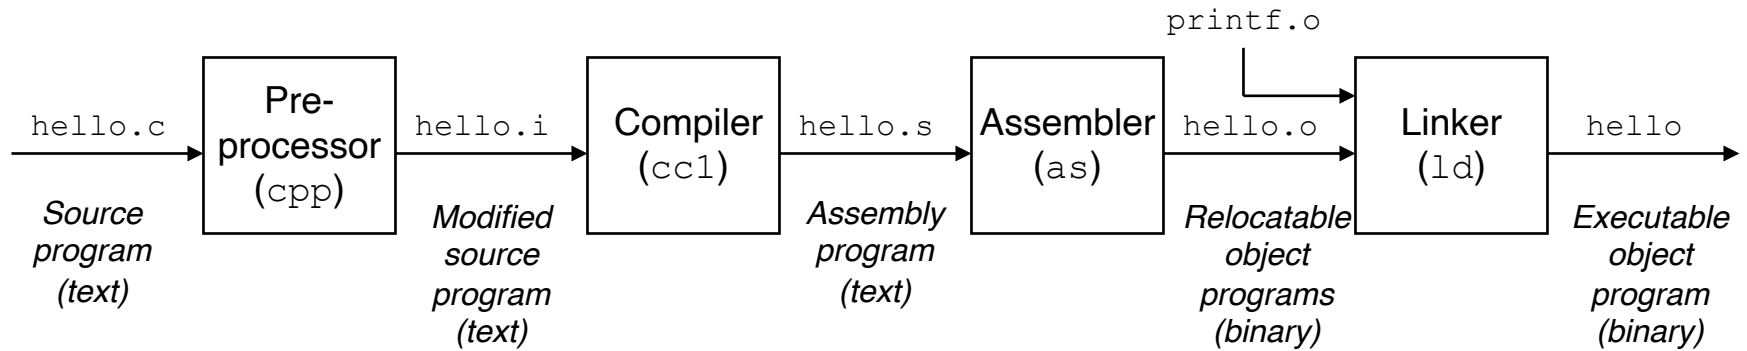
Figure 1.2 The ASCII text representation of hello.c.

# Programs translated by other programs

- `unix> gcc -o hello hello.c`



## Compilation System



- Pre-processing

- E.g., `#include<stdio.h>` is inserted into `hello.i`

- Compilation (.s)

- Each statement is an assembly language program

- Assembly (.o)

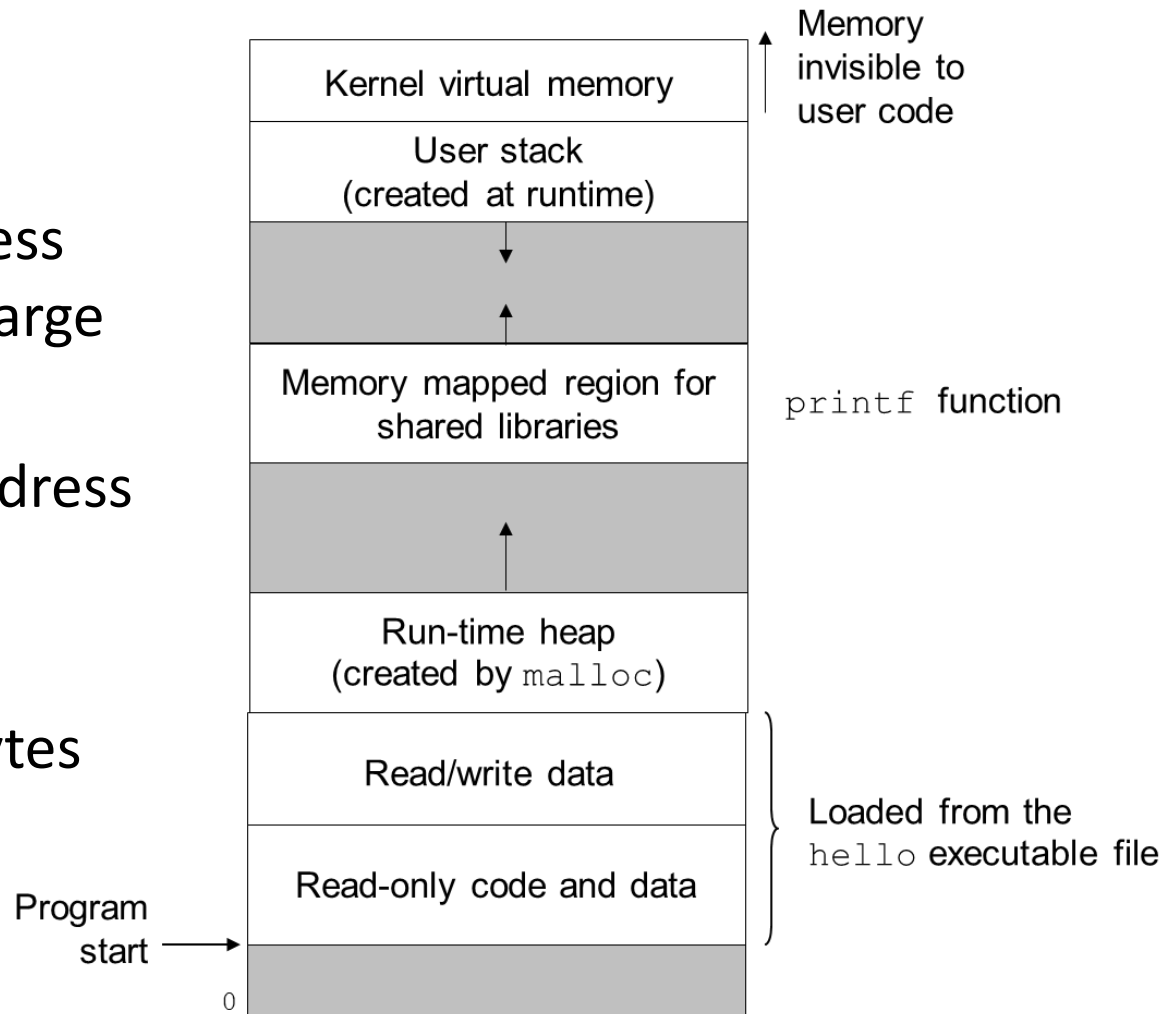
- A binary file whose bytes encode mach. language instructions

- Linking

- Get `printf()` which resides in a separate precompiled object file

# Virtual memory

- Illusion that each process has exclusive use of a large main memory
  - Example: Virtual address space for Linux
- **Files:** A sequence of bytes





# Amdahl's Law

- Effectiveness of improving the performance of one part of system
- Speed up one part → Effect on the overall system performance?
- $T_{new} = (1 - \alpha)T_{old} + \frac{\alpha T_{old}}{k}$
- $= T_{old} [(1 - \alpha) + \frac{\alpha}{k}]$
- $S = \frac{T_{old}}{T_{new}}$
- $S = \frac{1}{((1 - \alpha) + (\alpha/k))}$

# Amdahl's Law / Example

- Consider a system;
  - A part of the system initially consumed 60% of the time ( $\alpha = 0.6$ )
  - It is sped up by a factor of **3** ( $k=3$ )
- Overall improvement ?

# Amdahl's Law / Example

- Consider a system;
  - A part of the system initially consumed 60% of the time ( $\alpha = 0.6$ )
  - It is sped up by a factor of **3** ( $k=3$ )
- Overall improvement ?
- $$S = \frac{1}{((1-\alpha)+(\alpha/k))}$$
- $= 1 / [0.4 + 0.6/3] = \mathbf{1.67 \text{ times}}$

# Amdahl's Law / Example 2

- Calculate the following improvements on a current system and decide which one is better
- **1)** if we make 90% of a program run 10 times faster.
- **2)** if we make 80% of a program run 20% faster

# Amdahl's Law / Example 2

- **1)** if we make 90% of a program run 10 times faster.

$$S = \frac{1}{((1-\alpha)+(\alpha/k))} = \frac{1}{((1-0.9)+(0.9/10))} = \mathbf{5.26}$$

- **2)** if we make 80% of a program run 20% faster

$$S = \frac{1}{((1-\alpha)+(\alpha/k))} = \frac{1}{((1-0.8)+(0.8/1.2))} = \mathbf{1.153}$$

## Data Measurement Chart

Data Measurement	Size
------------------	------

Bit	Single Binary Digit (1 or 0)
-----	---------------------------------

Byte	8 bits
------	--------

Kilobyte (KB)	1,024 Bytes
---------------	-------------

Megabyte (MB)	1,024 Kilobytes
---------------	-----------------

Gigabyte (GB)	1,024 Megabytes
---------------	-----------------

Terabyte (TB)	1,024 Gigabytes
---------------	-----------------

Petabyte (PB)	1,024 Terabytes
---------------	-----------------

Exabyte (EB)	1,024 Petabytes
--------------	-----------------

# The Decimal System

- System based on decimal digits (0, 1, 2, 3, 4, 5, 6, 7, 8, 9) to represent numbers

- For example the number 83 means eight tens plus three:

$$83 = (8 * 10) + 3$$

- The number 4728 means four thousands, seven hundreds, two tens, plus eight:

$$4728 = (4 * 1000) + (7 * 100) + (2 * 10) + 8$$

- The decimal system is said to have a **base**, or **radix**, of 10. This means that each digit in the number is multiplied by 10 raised to a power corresponding to that digit's position:

$$83 = (8 * 10^1) + (3 * 10^0)$$

$$4728 = (4 * 10^3) + (7 * 10^2) + (2 * 10^1) + (8 * 10^0)$$

# The Binary System

- Only two digits, 1 and 0
- Represented to the base 2
- The digits 1 and 0 in binary notation have the same meaning as in decimal notation:

$$0_2 = 0_{10}$$

$$1_2 = 1_{10}$$

- To represent larger numbers each digit in a binary number has a value depending on its position:

$$10_2 = (1 * 2^1) + (0 * 2^0) = 2_{10}$$

$$11_2 = (1 * 2^1) + (1 * 2^0) = 3_{10}$$

$$100_2 = (1 * 2^2) + (0 * 2^1) + (0 * 2^0) = 4_{10}$$



For representing numbers in base 2, there are two possible digits (0, 1) in which column values are a power of two:

$2^7$	$2^6$	$2^5$	$2^4$	$2^3$	$2^2$	$2^1$	$2^0$
128	64	32	16	8	4	2	1
<hr/>							
0	1	1	0	0	0	1	1
0 +	64 +	32 +	0 +	0 +	0 +	2 +	1 = 99

Although values represented in base 2 are significantly longer than those in base 10, **binary representation is used in digital computing because of the resulting simplicity of hardware design**

# Encoding Byte Values

- Byte = 8 bits
  - Binary  $00000000_2$  to  $11111111_2$
  - Decimal:  $0_{10}$  to  $255_{10}$
  - Hexadecimal  $00_{16}$  to  $FF_{16}$ 
    - Base 16 number representation
    - Use characters '0' to '9' and 'A' to 'F'
    - Write  $FA1D37B_{16}$  in C as
      - `0xFA1D37B`
      - `0xfa1d37b`

Hex	Decimal	Binary
0	0	0000
1	1	0001
2	2	0010
3	3	0011
4	4	0100
5	5	0101
6	6	0110
7	7	0111
8	8	1000
9	9	1001
A	10	1010
B	11	1011
C	12	1100
D	13	1101
E	14	1110
F	15	1111

1100 1001 0111 1011 → 0xC97B

# Machine Words

- Machine has “word size”
  - Nominal size of integer-valued data
  - More importantly – a virtual address is encoded by such a word
    - Hence, it determines max size of virtual address space
  - Most current machines are 32 bits (4 bytes)
    - Limits addresses to 4GB
    - Becoming too small for memory-intensive applications
  - Newer systems are 64 bits (8 bytes)
    - Potentially address  $\approx 1.84 \times 10^{19}$  bytes
  - Machines support multiple data formats
    - Fractions or multiples of word size
    - Always integral number of bytes

# Data Sizes

- Each computer has a word size
  - For a machine with  $w$ -bit word size
  - The virtual address can range from 0 to  $2^w - 1$
  - The program access to at most  $2^w$  bytes
- 32 bit vs 64 bit


# Addressing and Byte Ordering

- For objects that span multiple bytes (e.g. integers), we need to agree on two things
  - what would be the address of the object?
  - how would we order the bytes in memory?


# Byte Ordering

- How to order bytes within multi-byte word in memory
- Conventions
  - (most) Sun's, IBMs are “Big Endian” machines
    - Least significant byte has highest address (comes last)
  - (most) Intel's are “Little Endian” machines
    - Least significant byte has lowest address (comes first)
- Example
  - Variable x has 4-byte representation 0x01234567
  - Address given by &x is 0x100 0x100 0x101

## Big Endian



## Little Endian



# Boolean Variables and Operations

- Developed by George Boole in 19th Century
  - Algebraic representation of logic
    - Encode “True” as 1 and “False” as 0
  - $\langle \{0,1\}, |, \&, \sim, 0, 1 \rangle$
  - $|$  is “sum” operation,  $\&$  is “product” operation
  - $\sim$  is “complement” operation (not additive inverse)
  - 0 is identity for sum, 1 is identity for product
- Makes use of variables and operations
  - Are logical
  - A variable may take on the value 1 (TRUE) or 0 (FALSE)
  - Basic logical operations are AND, OR, XOR and NOT

# Boolean Variables and Operations / 2

- AND

- Yields true (binary value 1) if and only if both of its operands are true
- In the absence of parentheses the AND operation takes precedence over the OR operation
- When no ambiguity will occur the AND operation is represented by simple concatenation instead of the dot operator

- OR

- Yields true if either or both of its operands are true

- NOT

- Inverts the value of its operand



# Table: Boolean Operators

(a) Boolean Operators of Two Input Variables

P	Q	NOT P ( $\bar{P}$ )	P AND Q ( $P \cdot Q$ )	P OR Q ( $P + Q$ )	P NAND Q ( $\overline{P \cdot Q}$ )	P NOR Q ( $\overline{P + Q}$ )	P XOR Q ( $P \oplus Q$ )
0	0	1	0	0	1	1	0
0	1	1	0	1	1	0	1
1	0	0	0	1	1	0	1
1	1	0	1	1	0	0	0

(b) Boolean Operators Extended to More than Two Inputs (A, B, ...)

Operation	Expression	Output = 1 if
AND	$A \cdot B \cdot \dots$	All of the set {A, B, ...} are 1.
OR	$A + B + \dots$	Any of the set {A, B, ...} are 1.
NAND	$\overline{A \cdot B \cdot \dots}$	Any of the set {A, B, ...} are 0.
NOR	$\overline{A + B + \dots}$	All of the set {A, B, ...} are 0.
XOR	$A \oplus B \oplus \dots$	The set {A, B, ...} contains an odd number of ones.

# Table: Basic Identities of Boolean Algebra

## Basic Postulates

$$A \cdot B = B \cdot A$$

$$A + B = B + A$$

Commutative Laws

$$A \cdot (B + C) = (A \cdot B) + (A \cdot C)$$

$$A + (B \cdot C) = (A + B) \cdot (A + C)$$

Distributive Laws

$$1 \cdot A = A$$

$$0 + A = A$$

Identity Elements

$$A \cdot \bar{A} = 0$$

$$A + \bar{A} = 1$$

Inverse Elements

## Other Identities

$$0 \cdot A = 0$$

$$1 + A = 1$$

$$A \cdot A = A$$

$$A + A = A$$

$$A \cdot (B \cdot C) = (A \cdot B) \cdot C$$

$$A + (B + C) = (A + B) + C$$

Associative Laws

$$\overline{A \cdot B} = \bar{A} + \bar{B}$$

$$\overline{A + B} = \bar{A} \cdot \bar{B}$$

DeMorgan's Theorem

# Exercise 1


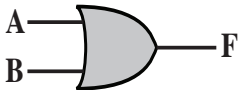
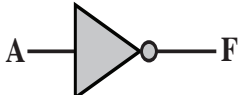

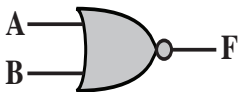
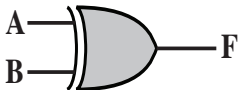
Evaluate the following expression when  $A = 0$ ,  $B = 1$ , and  $C = 1$

$$F = B + \bar{C}A + B\bar{A} + A\bar{B}$$

Simplify the following functions;

$$F = AB + BC + \bar{B}C$$

$$F = A + \bar{A}B$$

Name	Graphical Symbol	Algebraic Function	Truth Table															
AND		$F = A \cdot B$ or $F = AB$	<table><tr><th>A</th><th>B</th><th>F</th></tr><tr><td>0</td><td>0</td><td>0</td></tr><tr><td>0</td><td>1</td><td>0</td></tr><tr><td>1</td><td>0</td><td>0</td></tr><tr><td>1</td><td>1</td><td>1</td></tr></table>	A	B	F	0	0	0	0	1	0	1	0	0	1	1	1
A	B	F																
0	0	0																
0	1	0																
1	0	0																
1	1	1																
OR		$F = A + B$	<table><tr><th>A</th><th>B</th><th>F</th></tr><tr><td>0</td><td>0</td><td>0</td></tr><tr><td>0</td><td>1</td><td>1</td></tr><tr><td>1</td><td>0</td><td>1</td></tr><tr><td>1</td><td>1</td><td>1</td></tr></table>	A	B	F	0	0	0	0	1	1	1	0	1	1	1	1
A	B	F																
0	0	0																
0	1	1																
1	0	1																
1	1	1																
NOT		$F = \bar{A}$ or $F = A'$	<table><tr><th>A</th><th>F</th></tr><tr><td>0</td><td>1</td></tr><tr><td>1</td><td>0</td></tr></table>	A	F	0	1	1	0									
A	F																	
0	1																	
1	0																	
NAND		$F = \overline{AB}$	<table><tr><th>A</th><th>B</th><th>F</th></tr><tr><td>0</td><td>0</td><td>1</td></tr><tr><td>0</td><td>1</td><td>1</td></tr><tr><td>1</td><td>0</td><td>1</td></tr><tr><td>1</td><td>1</td><td>0</td></tr></table>	A	B	F	0	0	1	0	1	1	1	0	1	1	1	0
A	B	F																
0	0	1																
0	1	1																
1	0	1																
1	1	0																
NOR		$F = \overline{A + B}$	<table><tr><th>A</th><th>B</th><th>F</th></tr><tr><td>0</td><td>0</td><td>1</td></tr><tr><td>0</td><td>1</td><td>0</td></tr><tr><td>1</td><td>0</td><td>0</td></tr><tr><td>1</td><td>1</td><td>0</td></tr></table>	A	B	F	0	0	1	0	1	0	1	0	0	1	1	0
A	B	F																
0	0	1																
0	1	0																
1	0	0																
1	1	0																
XOR		$F = A \oplus B$	<table><tr><th>A</th><th>B</th><th>F</th></tr><tr><td>0</td><td>0</td><td>0</td></tr><tr><td>0</td><td>1</td><td>1</td></tr><tr><td>1</td><td>0</td><td>1</td></tr><tr><td>1</td><td>1</td><td>0</td></tr></table>	A	B	F	0	0	0	0	1	1	1	0	1	1	1	0
A	B	F																
0	0	0																
0	1	1																
1	0	1																
1	1	0																

# General Boolean Algebras

- Boolean operations can be extended to work on bit vectors
  - Operations applied bitwise

01101001	01101001	01101001	
& 01010101	01010101	^ 01010101	~ 01010101
01000001	01111101	00111100	10101010

- All of the properties of Boolean algebra apply
- Now, Boolean |, & and ~ correspond to set union, intersection and complement

# Bit-Level Operations in C

- Operations **&**, **|**, **~**, **^** Available in C
  - Apply to any “integral” data type
    - long, int, short, char, unsigned
  - View arguments as bit vectors
  - Arguments applied bit-wise
- Examples (Char data type)
  - $\sim 0x41 \rightarrow 0xBE$ 
    - $\sim 01000001_2 \rightarrow 10111110_2$
  - $\sim 0x00 \rightarrow 0xFF$ 
    - $\sim 00000000_2 \rightarrow 11111111_2$
  - $0x69 \& 0x55 \rightarrow 0x41$ 
    - $01101001_2 \& 01010101_2 \rightarrow 01000001_2$
  - $0x69 | 0x55 \rightarrow 0x7D$ 
    - $01101001_2 | 01010101_2 \rightarrow 01111101_2$

# Contrast: Logic Operations in C

- Contrast to Logical Operators

– **&&, ||, !**

- View 0 as “False”
- Anything nonzero as “True”
- Always return 0 or 1
- **Early termination**

- Examples (char data type)

– !0x41 → 0x00  
– !0x00 → 0x01  
– !!0x41 → 0x01

– 0x69 && 0x55 → 0x01  
– 0x69 || 0x55 → 0x01

– p && \*p (avoids null pointer access)

Watch out for && vs. & (and || vs. |) ...  
one of the more common oopsies  
in C programming

# Shift Operations

- Left Shift:  $x \ll y$ 
  - Shift bit-vector **x** left **y** positions
    - Throw away extra bits on left
      - Fill with 0's on right
- Right Shift:  $x \gg y$ 
  - Shift bit-vector **x** right **y** positions
    - Throw away extra bits on right
  - Logical shift
    - Fill with 0's on left
  - Arithmetic shift
    - Replicate most significant bit on left
    - Useful in two's complement
- Undefined Behavior
  - Shift amount  $< 0$  or  $\geq$  word size

Argument <b>x</b>	01100010
$\ll 3$	00010000
Log. $\gg 2$	00011000
Arith. $\gg 2$	00011000

Argument <b>x</b>	10100010
$\ll 3$	00010000
Log. $\gg 2$	00101000
Arith. $\gg 2$	11101000



# Boolean Algebra $\approx$ Integer Ring

Commutative	$A \mid B = B \mid A$ $A \& B = B \& A$	$A + B = B + A$ $A * B = B * A$
Associativity	$(A \mid B) \mid C = A \mid (B \mid C)$ $(A \& B) \& C = A \& (B \& C)$	$(A + B) + C = A + (B + C)$ $(A * B) * C = A * (B * C)$
Product distributes over sum	$A \& (B \mid C) = (A \& B) \mid (A \& C)$	$A * (B + C) = A * B + B * C$
Sum and product identities	$A \mid 0 = A$ $A \& 1 = A$	$A + 0 = A$ $A * 1 = A$
Zero is product annihilator	$A \& 0 = 0$	$A * 0 = 0$
Cancellation of negation	$\sim (\sim A) = A$	$- (-A) = A$

# Boolean Algebra $\neq$ Integer Ring

Boolean: Sum distributes over product	$A \mid (B \ \& \ C) = (A \mid B) \ \& \ (A \mid C)$	$A + (B * C) \neq (A + B) * (B + C)$
Boolean: Idempotency	$A \mid A = A$ $A \ \& \ A = A$	$A + A \neq A$ $A * A \neq A$
Boolean: Absorption	$A \mid (A \ \& \ B) = A$ $A \ \& \ (A \mid B) = A$	$A + (A * B) \neq A$ $A * (A + B) \neq A$
Boolean: Laws of Complements	$A \mid \sim A = 1$	$A + -A \neq 1$
Ring: Every element has additive inverse	$A \mid \sim A \neq 0$	$A + -A = 0$

# Properties of & and ^

- Boolean ring
  - $\langle \{0,1\}, ^, \&, I, 0, 1 \rangle$
  - Identical to integers mod 2
  - I is identity operation:  $I(A) = A$ 
    - $A \wedge A = 0$

- Property: Boolean ring
  - Commutative sum
  - Commutative product
  - Associative sum
  - Associative product
  - Prod. over sum
  - 0 is sum identity
  - 1 is prod. identity
  - 0 is product annihilator
  - Additive inverse

$$A \wedge B = B \wedge A$$

$$A \& B = B \& A$$

$$(A \wedge B) \wedge C = A \wedge (B \wedge C)$$

$$(A \& B) \& C = A \& (B \& C)$$

$$A \& (B \wedge C) = (A \& B) \wedge (A \& C)$$

$$A \wedge 0 = 0$$

$$A \& 1 = A$$

$$A \& 0 = 0$$

$$A \wedge A = 0$$

# Unsigned Encodings

$$B2U(X) = \sum_{i=0}^{w-1} x_i \cdot 2^i$$

(Binary To Unsigned)


e.g.  $B2U([1011]) = 1 * 2^3 + 0 * 2^2 + 1 * 2^1 + 1 * 2^0 = 11$

– C short 2 bytes long

```
short int x = 15213;
```

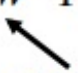
	Decimal	Hex	Binary
x	15213	3B 6D	00111011 01101101

# Examples



# Two's Complement

$$B2T(X) = -x_{w-1} \cdot 2^{w-1} + \sum_{i=0}^{w-2} x_i \cdot 2^i$$

  
**Sign Bit**

– e.g.  $B2T([1011]) = -1 * 2^3 + 0 * 2^2 + 1 * 2^1 + 1 * 2^0 = -5$

– C `short` 2 bytes long

```
short int y = -15213;
```

	Decimal	Hex	Binary
x	15213	3B 6D	00111011 01101101
y	-15213	C4 93	11000100 10010011

- **Sign bit**

- For 2's complement, most significant bit indicates sign
  - 0 for nonnegative; 1 for negative

# Two's Complement

- Invert and add **one**

Suppose we're working with 8 bit quantities and suppose we want to find how **-28** would be expressed in two's complement notation.

- First we write out 28 in **binary form**.

00011100

- Then we **invert the digits**. 0 becomes 1, 1 becomes 0.

11100011

- Then we **add 1**.

11100100

That is how one would write -28 in 8 bit binary.

# Numeric Ranges

## ■ Unsigned Values

- $UMin = 0$   
000...0

- $UMax = 2^w - 1$   
111...1

## ■ Two's Complement Values

- $TMin = -2^{w-1}$   
100...0

- $TMax = 2^{w-1} - 1$   
011...1

## ■ Other Values

- Minus 1  
111...1

### Values for $W = 16$

	Decimal	Hex	Binary
<b>UMax</b>	<b>65535</b>	<b>FF FF</b>	<b>11111111 11111111</b>
<b>TMax</b>	<b>32767</b>	<b>7F FF</b>	<b>01111111 11111111</b>
<b>TMin</b>	<b>-32768</b>	<b>80 00</b>	<b>10000000 00000000</b>
<b>-1</b>	<b>-1</b>	<b>FF FF</b>	<b>11111111 11111111</b>
<b>0</b>	<b>0</b>	<b>00 00</b>	<b>00000000 00000000</b>



# Signed vs. Unsigned in C

- Constants

- By default are considered to be signed integers
- Unsigned if have “U” as suffix

`0U, 4294967259U`

- Casting

- Explicit casting between signed & unsigned same as U2T and T2U


```
int tx, ty;  
unsigned ux, uy;  
tx = (int) ux;  
uy = (unsigned) ty;
```

- Implicit casting also occurs via assignments and procedure calls

```
tx = ux;  
uy = ty;
```

# Sign Extension

- **Task:**
  - Given  $w$ -bit signed integer  $x$
  - Convert it to  $w+k$ -bit integer with same value
- **Rule:**
  - Make  $k$  copies of sign bit:
  - $X' = x_{w-1}, \dots, x_{w-1}, x_{w-1}, x_{w-2}, \dots, x_0$



# Truncating Numbers

- Reduce the number of bits representing the number
- Truncating  $w$ -bit number to a  $k$  bit number, we drop the high order  $w-k$  bits
  - Can alter its value
    - A form of overflow

# Summary: Expanding, Truncating: Basic Rules

- **Expanding (e.g., short int to int)**
  - Unsigned: zeros added
  - Signed: sign extension
  - Both yield expected result
- **Truncating (e.g., unsigned to unsigned short)**
  - Unsigned/signed: bits are truncated
  - Result reinterpreted
  - Unsigned: mod operation
  - Signed: similar to mod
  - For small numbers yields expected behavior

# OVERFLOW RULE:

- If two numbers are added, and they are both positive or both negative, then overflow occurs if and only if the result has the opposite sign.

# Multiplication

- Goal: Computing Product of  $w$ -bit numbers  $x, y$ 
  - Either signed or unsigned
- But, exact results can be bigger than  $w$  bits
  - Unsigned: up to  $2w$  bits
    - Result range:  $0 \leq x * y \leq (2^w - 1)^2 = 2^{2w} - 2^{w+1} + 1$
  - Two's complement min (negative): Up to  $2w-1$  bits
    - Result range:  $x * y \geq (-2^{w-1}) * (2^{w-1} - 1) = -2^{2w-2} + 2^{w-1}$
  - Two's complement max (positive): Up to  $2w$  bits, but only for  $(TMin_w)^2$ 
    - Result range:  $x * y \leq (-2^{w-1})^2 = 2^{2w-2}$
- So, maintaining exact results...
  - would need to keep expanding word size with each product computed
  - is done in software, if needed
    - e.g., by “arbitrary precision” arithmetic packages

# Unsigned Binary Multiplication

1011	<b>Multiplicand (11)</b>
× 1101	<b>Multiplier (13)</b>
-----	
1011	} <b>Partial products</b>
0000	
1011	
1011	
-----	
10001111	<b>Product (143)</b>


**Figure 10.7** Multiplication of Unsigned Binary Integers

# Power-of-2 Multiply with Shift

## ■ Operation

- $u \ll k$  gives  $u * 2^k$
- Both signed and unsigned

Operands:  $w$  bits



True Product:  $w+k$  bits      $u \cdot 2^k$

Discard  $k$  bits:  $w$  bits

UMult <sub>$w$</sub> ( $u, 2^k$ )  
TMult <sub>$w$</sub> ( $u, 2^k$ )

## ■ Examples


- $u \ll 3 \quad \quad \quad == \quad u * 8$
- $(u \ll 5) - (u \ll 3) == u * 24$
- Most machines shift and add faster than multiply
  - Compiler generates this code automatically



# Unsigned Power-of-2 Divide with Shift

## ■ Quotient of Unsigned by Power of 2

- $u \gg k$  gives  $\lfloor u / 2^k \rfloor$
- Uses logical shift



	Division	Computed	Hex	Binary
<b>x</b>	<b>15213</b>	<b>15213</b>	3B 6D	00111011 01101101
<b>x &gt;&gt; 1</b>	<b>7606.5</b>	<b>7606</b>	1D B6	00011101 10110110
<b>x &gt;&gt; 4</b>	<b>950.8125</b>	<b>950</b>	03 B6	00000011 10110110
<b>x &gt;&gt; 8</b>	<b>59.4257813</b>	<b>59</b>	00 3B	00000000 00111011

# Arithmetic: Basic Rules

## ■ Addition:

- Unsigned/signed: Normal addition followed by truncate, same operation on bit level
- Unsigned: addition mod  $2^w$ 
  - Mathematical addition + possible subtraction of  $2^w$
- Signed: modified addition mod  $2^w$  (result in proper range)
  - Mathematical addition + possible addition or subtraction of  $2^w$

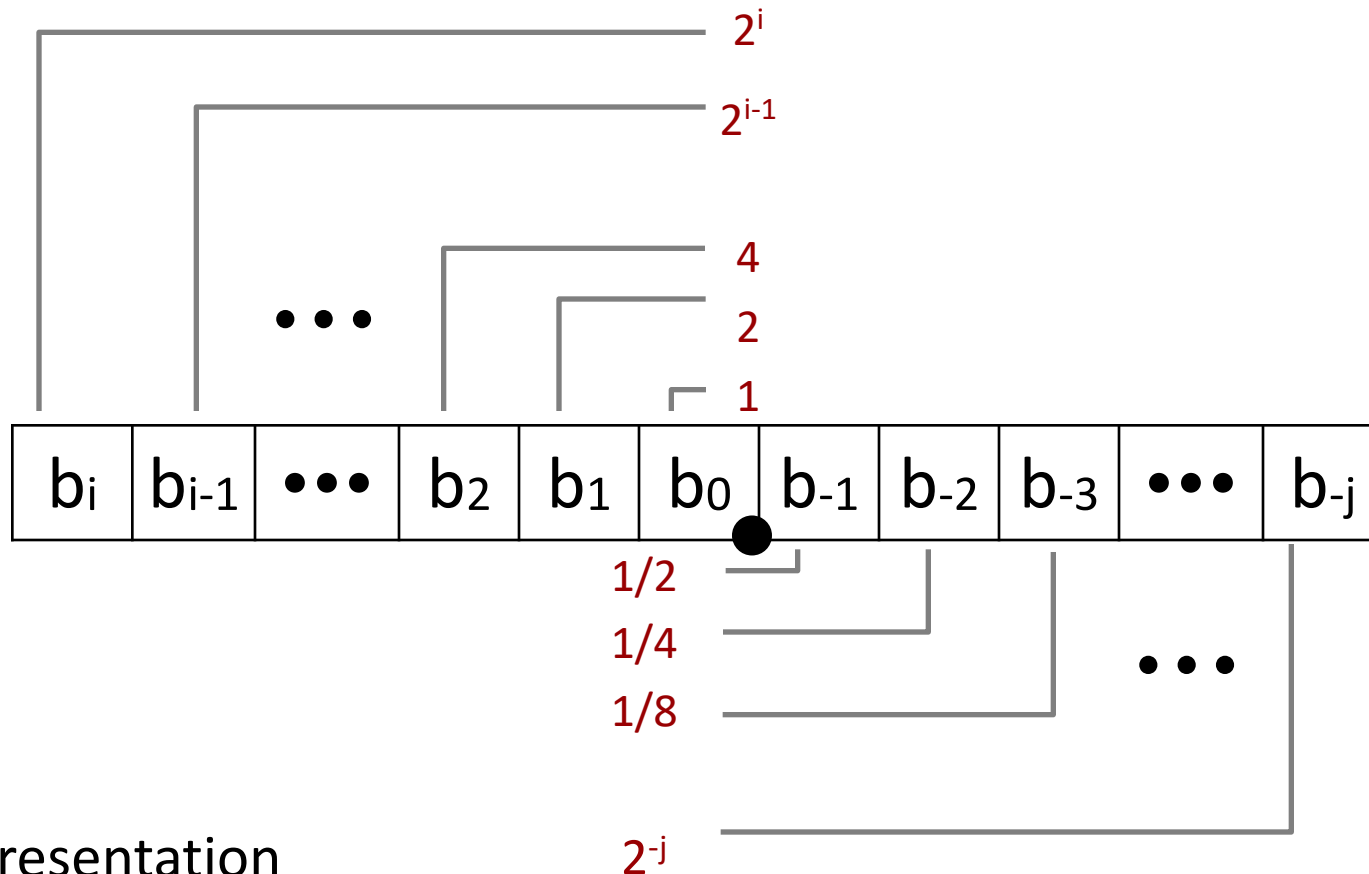
## ■ Multiplication:

- Unsigned/signed: Normal multiplication followed by truncate, same operation on bit level
- Unsigned: multiplication mod  $2^w$
- Signed: modified multiplication mod  $2^w$  (result in proper range)

# Fractional binary numbers

- What is  $1011.101_2$ ?

# Fractional Binary Numbers



- Representation

- Bits to right of “binary point” represent fractional powers of 2
- Represents rational number:

$$\sum_{k=-j}^i b_k \times 2^k$$

# Fractional Binary Numbers: Examples

## ■ Value Representation

$5 \frac{3}{4}$	$101.11_2$
$2 \frac{7}{8}$	$10.111_2$
$1 \frac{7}{16}$	$1.0111_2$

## ■ Observations

- Divide by 2 by shifting right (unsigned)
- Multiply by 2 by shifting left
- Numbers of form  $0.111111..._2$  are just below 1.0
  - $1/2 + 1/4 + 1/8 + \dots + 1/2^i + \dots \rightarrow 1.0$
  - Use notation  $1.0 - \epsilon$

# Representable Numbers

- **Limitation #1**

- Can only exactly represent numbers of the form  $x/2^k$ 
  - Other rational numbers have repeating bit representations

- **Value      Representation**

- $1/3$        $0.0101010101 [01] \dots_2$
- $1/5$        $0.001100110011 [0011] \dots_2$
- $1/10$        $0.0001100110011 [0011] \dots_2$

- **Limitation #2**

- Just one setting of binary point within the  $w$  bits
  - Limited range of numbers (very small values? very large?)

# IEEE Floating Point

- **IEEE Standard 754**
  - Established in 1985 as uniform standard for floating point arithmetic
    - Before that, many idiosyncratic formats
  - Supported by all major CPUs
- **Driven by numerical concerns**
  - Nice standards for rounding, overflow, underflow
  - Hard to make fast in hardware
    - Numerical analysts predominated over hardware designers in defining standard

# Floating Point Representation

- Numerical Form:

$$(-1)^s M 2^E$$

- Sign bit **s** determines whether number is negative or positive
- Significand **M** normally a fractional value in range [1.0,2.0).
- Exponent **E** weights value by power of two

- Encoding

- MSB **S** is sign bit **s**
- exp field encodes **E** (but is not equal to E)
- frac field encodes **M** (but is not equal to M)





# Precision options

- Single precision: 32 bits



- Double precision: 64 bits



- The value encoded by a given bit representation can be divided into three different cases, depending on the value of **exp**.
- Case 1: Normalized Values
- Case 2: Denormalized Values
- Case 3: Special Values

# Case 1: “Normalized” Values

$$v = (-1)^s M 2^E$$

- When:  $exp \neq 000...0$  and  $exp \neq 111...1$
- Exponent coded as a biased value:  $E = \text{Exp} - \text{Bias}$ 
  - **Exp**: unsigned value of exp field
  - **Bias** =  $2^{k-1} - 1$ , where k is number of exponent bits
    - Single precision: 127 (Exp: 1...254, E: -126...127)
    - Double precision: 1023 (Exp: 1...2046, E: -1022...1023)
- Significand coded with implied leading 1:  $M = 1.\text{xxx}...\text{x}_2$ 
  - xxx...x: bits of frac field
  - Minimum when frac=000...0 ( $M = 1.0$ )
  - Maximum when frac=111...1 ( $M = 2.0 - \epsilon$ )
  - Get extra leading bit for “free”

# Case 2: Denormalized Values

$$v = (-1)^s M 2^E$$
$$E = 1 - \text{Bias}$$

- **Condition:  $\text{exp} = 000\dots 0$**
- Exponent value:  $E = 1 - \text{Bias}$  (instead of  $E = 0 - \text{Bias}$ )
- Significand coded with implied leading 0:  $M = 0.\text{xxx}\dots\text{x}_2$ 
  - **$\text{xxx}\dots\text{x}$** : bits of **frac**
- Cases
  - **$\text{exp} = 000\dots 0, \text{frac} = 000\dots 0$** 
    - Represents zero value
    - Note distinct values:  $+0$  and  $-0$  (why?)
  - **$\text{exp} = 000\dots 0, \text{frac} \neq 000\dots 0$** 
    - Numbers closest to  $0.0$
    - Equispaced

# Case 3: Special Values

- Condition: **exp** = 111...1
- Case: **exp** = 111...1, **frac** = 000...0
  - Represents value  $\infty$  (infinity)
  - Operation that overflows
  - Both positive and negative
  - E.g.,  $1.0/0.0 = -1.0/-0.0 = +\infty$ ,  $1.0/-0.0 = -\infty$
- Case: **exp** = 111...1, **frac**  $\neq$  000...0
  - Not-a-Number (NaN)
  - Represents case when no numeric value can be determined
  - E.g.,  $\text{sqrt}(-1)$ ,  $\infty - \infty$ ,  $\infty \times 0$

# Special Properties of the IEEE Encoding

- **FP Zero Same as Integer Zero**
  - All bits = 0
- **Can (Almost) Use Unsigned Integer Comparison**
  - Must first compare sign bits
  - Must consider  $-0 = 0$
  - NaNs problematic
    - Will be greater than any other values
    - What should comparison yield?
  - Otherwise OK
    - Denorm vs. normalized
    - Normalized vs. infinity

# Floating Point Operations: Basic Idea

- $x +_f y = \text{Round}(x + y)$
- $x \times_f y = \text{Round}(x \times y)$
- **Basic idea**
  - First **compute exact result**
  - Make it fit into desired precision
    - Possibly overflow if exponent too large
    - Possibly **round to fit into frac**

# Rounding

- Rounding Modes (illustrate with \$ rounding)

•	<b>\$1.40</b>	<b>\$1.60</b>	<b>\$1.50</b>	<b>\$2.50</b>	<b>−\$1.50</b>
– Towards zero	\$1	\$1	\$1	\$2	−\$1
– Round down ( $-\infty$ )	\$1	\$1	\$1	\$2	−\$2
– Round up ( $+\infty$ )	\$2	\$2	\$2	\$3	−\$1
– Nearest Even (default)	\$1	\$2	\$2	\$2	−\$2



# FP Multiplication

- $(-1)^{s1} M1 2^{E1} \times (-1)^{s2} M2 2^{E2}$
- **Exact Result:**  $(-1)^s M 2^E$ 
  - Sign  $s$ :  $s1 \wedge s2$
  - Significand  $M$ :  $M1 \times M2$
  - Exponent  $E$ :  $E1 + E2$
- **Fixing**
  - If  $M \geq 2$ , shift  $M$  right, increment  $E$
  - If  $E$  out of range, overflow
  - Round  $M$  to fit **frac** precision
- **Implementation**
  - Biggest chore is multiplying significands

# Floating Point Addition

- $(-1)^{s1} M1 2^{E1} + (-1)^{s2} M2 2^{E2}$

—Assume  $E1 > E2$

- **Exact Result:**  $(-1)^s M 2^E$

—Sign  $s$ , significand  $M$ :

- Result of signed align & add

—Exponent  $E$ :  $E1$

- **Fixing**


—If  $M \geq 2$ , shift  $M$  right, increment  $E$

—if  $M < 1$ , shift  $M$  left  $k$  positions, decrement  $E$  by  $k$

—Overflow if  $E$  out of range

—Round  $M$  to fit **frac** precision

Get binary points lined up



# Mathematical Properties of FP Add

- Compare to those of Abelian Group
  - Closed under addition? **Yes**
    - But may generate infinity or NaN
  - Commutative? **Yes**
  - Associative? **No**
    - Overflow and inexactness of rounding
    - $(3.14 + 1e10) - 1e10 = 0$ ,  $3.14 + (1e10 - 1e10) = 3.14$
  - 0 is additive identity? **Yes**
  - Every element has additive inverse? **Almost**
    - Yes, except for infinities & NaNs
- Monotonicity **Almost**
  - $a \geq b \Rightarrow a + c \geq b + c$ 
    - Except for infinities & NaNs

# Mathematical Properties of FP Mult

- **Compare to Commutative Ring**

- Closed under multiplication? Yes
  - But may generate infinity or NaN
- Multiplication Commutative? Yes
- Multiplication is Associative? No
  - Possibility of overflow, inexactness of rounding
  - Ex:  $(1e20 * 1e20) * 1e-20 = \text{inf}$ ,  $1e20 * (1e20 * 1e-20) = 1e20$
- 1 is multiplicative identity? Yes
- Multiplication distributes over addition? No
  - Possibility of overflow, inexactness of rounding
  - $1e20 * (1e20 - 1e20) = 0.0$ ,  $1e20 * 1e20 - 1e20 * 1e20 = \text{NaN}$

- **Monotonicity**

- $a \geq b \ \& \ c \geq 0 \Rightarrow a * c \geq b * c$  Almost
  - Except for infinities & NaNs

# Floating Point in C

- **C Guarantees Two Levels**
  - **float** single precision
  - **double** double precision
- **Conversions/Casting**
  - Casting between **int**, **float**, and **double** changes bit representation
  - **double/float** → **int**
    - Truncates fractional part
    - Like rounding toward zero
    - Not defined when out of range or NaN: Generally sets to TMin
  - **int** → **double**
    - Exact conversion, as long as **int** has  $\leq 53$  bit word size
  - **int** → **float**
    - Will round according to rounding mode

# Floating Point Puzzles

- For each of the following C expressions, either:
  - Argue that it is true for all argument values
  - Explain why not true

```
int x = ...;  
float f = ...;  
double d = ...;
```

Assume neither  
**d** nor **f** is NaN

- `x == (int)(float) x`
- `x == (int)(double) x`
- `f == (float)(double) f`
- `d == (double)(float) d`
- `f == -(-f);`
- `1.0/2 == 1/2.0`
- `d * d >= 0.0`
- `(d+f) - d == f`

# Intel x86 Processors

- Dominate laptop/desktop/server market
- Evolutionary design
  - Backwards compatible up until 8086, introduced in 1978
  - Added more features as time goes on
- Complex instruction set computer (CISC)
  - Many different instructions with many different formats
    - But, only small subset encountered with Linux programs
  - Hard to match performance of Reduced Instruction Set Computers (RISC)
  - But, Intel has done just that!
    - In terms of speed. Less so for low power.

# Our Coverage

- IA32
  - The traditional x86
  - For ???: RIP, Fall 2018
- x86-64
  - The standard
  - `moat.cis.uab.edu`
  - `gcc hello.c`
  - `gcc -m64 hello.c`
- Presentation
  - Book covers x86-64
  - So, we will cover x86-64

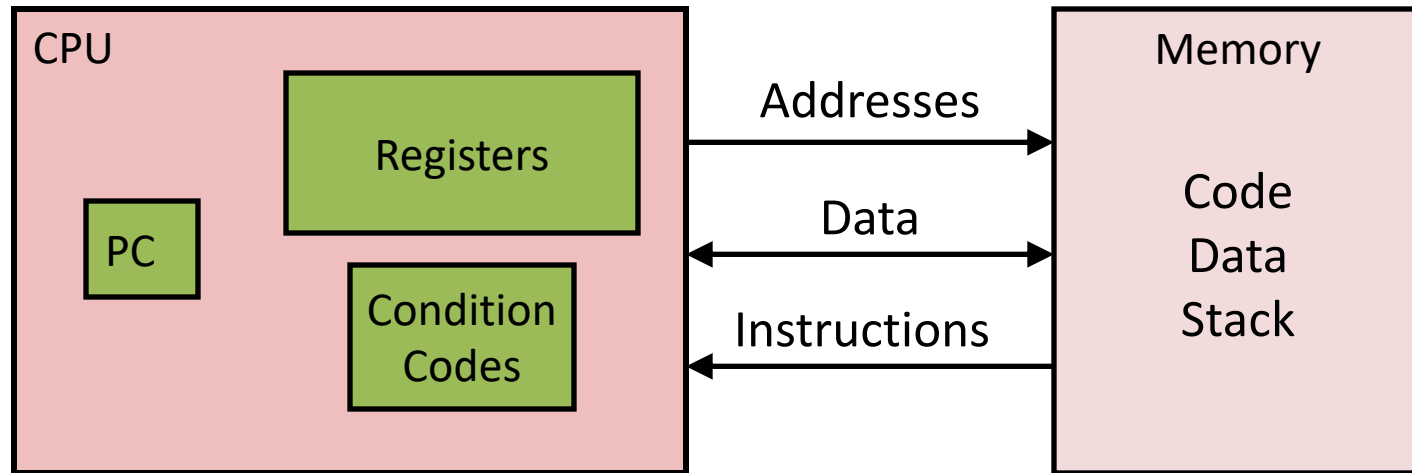


- **C, assembly, machine code**

## Definitions

- **Architecture:** (also ISA: instruction set architecture) The parts of a processor design that one needs to understand or write assembly/machine code.
  - Examples: instruction set specification, registers.
- **Microarchitecture:** Implementation of the architecture.
  - Examples: cache sizes and core frequency.
- Code Forms:
  - **Machine Code:** The byte-level programs that a processor executes
  - **Assembly Code:** A text representation of machine code
- Example ISAs:
  - Intel: x86, IA32, Itanium, x86-64
  - ARM: Used in almost all mobile phones

# Assembly/Machine Code View



## Programmer-Visible State

### – PC: Program counter

- Indicates the address of next instruction
- Called “%rip” (x86-64)

### – Register file

- Heavily used program data
- 16 named locations storing 64bit values

### – Condition codes


- Store status information about most recent arithmetic or logical operation
- Used for conditional branching

### – Memory

- Byte addressable array
- Code and user data
- Stack to support procedures

# Turning C into Object Code

- Code in files `p1.c` `p2.c`
- Compile with command: `gcc -Og p1.c p2.c -o p`
  - Use basic optimizations (`-Og`) [New to recent versions of GCC]
  - Put resulting binary in file `p`



# Compiling Into Assembly

## C Code (sum.c)

```
long plus(long x, long y);

void sumstore(long x, long y,
              long *dest)
{
    long t = plus(x, y);
    *dest = t;
}
```

## Generated x86-64 Assembly

```
sumstore:
    pushq    %rbx
    movq     %rdx, %rbx
    call     plus
    movq     %rax, (%rbx)
    popq     %rbx
    ret
```

Obtain (on shark machine) with command

```
gcc -Og -S sum.c
```

Produces file `sum.s`

**Warning:** This is the output of the textbook. We will get very different results on our machines due to different versions of gcc and different compiler settings.

# Vulcan server output

```
.file    "longplus.c"
.text
.globl   sumstore
.type    sumstore, @function
sumstore:
.LFB0:
        .cfi_startproc
        pushq   %rbx
        .cfi_def_cfa_offset 16
        .cfi_offset 3, -16
        movq    %rdx, %rbx
        call    plus
        movq    %rax, (%rbx)
        popq    %rbx
        .cfi_def_cfa_offset 8
        ret
        .cfi_endproc
.LFE0:
        .size    sumstore, .-sumstore
        .ident   "GCC: (GNU) 4.8.5 20150623 (Red Hat 4.8.5-28)"
        .section .note.GNU-stack,"",@progbits
```